

OpenSEO · one Claude Code session · Aug 18–19, 2026

Sprint Prompt Log

Every prompt from the overnight SEO sprint and what it produced. All voice-dictated — no prompt engineering, just talking to Claude Code like a colleague.

1. The overnight brief

"Use the SEO audit skill on our site. Do keyword research and build comparison pages for Semrush, Ahrefs, and Ubersuggest — not feature-for-feature: they're better for agencies, we're better if you work in Claude Code or \$100/mo makes you gasp. Come up with a backlink strategy, scrape emails, draft outreach. Run it overnight with adversarial reviewers."

One voice memo, three workflows, 26 agents. The two spoken sentences of positioning steered every writer.

MADE · [site audit](#) · [keyword research](#) · 3 comparison pages ([PR #504](#)) · [backlink strategy](#) + 13 verified outreach drafts · [skill learnings](#)

2. The follow-up question

"AI SEO tools is our strongest keyword — what content should we make to capture demand?"

One sentence, ~160 credits of live keyword + SERP data. Finding: "ai seo tools" is only the fourth-best keyword in its own cluster — "seo agent" and "aeo" are 27,100/mo at KD under 20.

MADE · [AI SEO Content Plan](#)

3. Artifacts + a second content wave

"Make Claude artifacts for all these outputs so I can walk through reviewing them. And go ahead and build the AI SEO content as another PR."

A 15-agent workflow wrote four pieces while the artifact suite was built in parallel.

MADE · [/seo-agent](#) · [AEO guide](#) · [tools roundup](#) · [SEO with Claude Code \(PR #506\)](#) · the [hub](#), [comparison review](#), and [backlink playbook](#)

4. The slide deck

"Make a minimal black-and-white slideshow: the three pillars of SEO. Technical — can your pages get found, do they load fast? Content — is it good, is it unique? Authority — do reputable sites link to you, do people talk about you?"

The dictated questions became the slides, verbatim.

MADE · [SEO Deck](#)

5. Design feedback, screenshot in hand

"De-slop these artifacts — super minimal, like our blog. Here's a screenshot." Then: "Remove the mustaches above and below the title. White only. Add a 'What is SEO?' slide."

Claude read the real design tokens out of the site's CSS and restyled the whole suite to match the blog. Mid-turn corrections landed while the work was running.

6. Making it reviewable

"Can I read the full research reports in a markdown viewer? And link new pages to localhost so I can click through."

Full reports published as rendered markdown; every page mention across the suite now opens the local dev server.

7. The bug report

"Can you fix the table styling?" (one screenshot: a dark table on the light page)

The screenshot was the whole report. Root cause found (OS dark mode flipping the site's prose-table tokens), fixed, committed to PR #506, verified on the dev server.

What the prompts have in common

- **Positioning in your own words, not specs** — two spoken sentences steered 26 agents.
- **Name the deliverable, not the steps** — the research → write → fact-check → copy-edit pipeline came from the agent.
- **Screenshots are prompts** — the restyle and the bug fix each took one.
- **Interrupt freely** — mid-turn corrections got folded in without restarting.
- **Ask for the meta-work** — "track what the agents struggle with" produced a skill-improvement report and six product bug candidates.